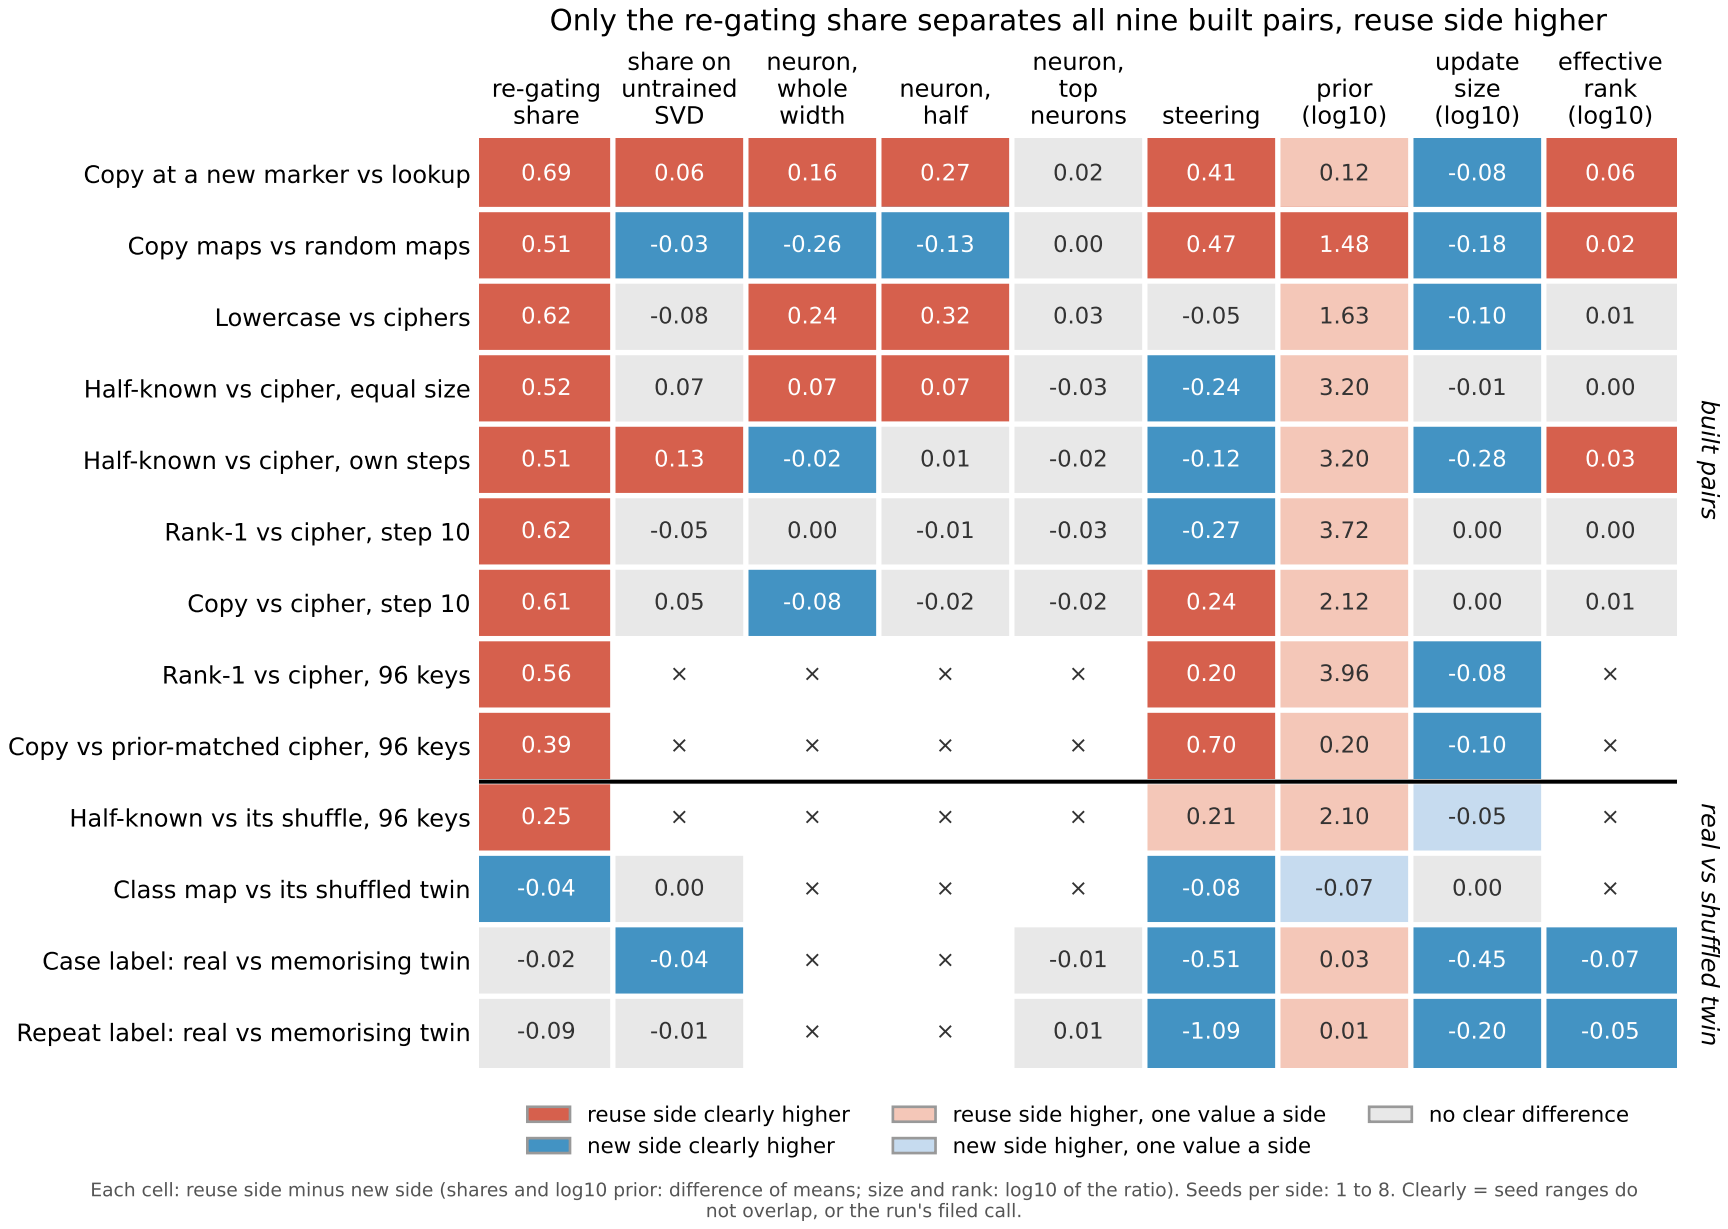
Only the re-gating share separates all nine built pairs, reuse side higher
re-gating share
share on untrained SVD
neuron, whole width
neuron, half
neuron, top neurons
steering
prior (log10)
update size (log10)
effective rank (log10)
Copy at a new marker vs lookup 0.69 0.06 0.16 0.27 0.02 0.41 0.12 -0.08 0.06
Copy maps vs random maps 0.51 -0.03 -0.26 -0.13 0.00 0.47 1.48 -0.18 0.02
Lowercase vs ciphers 0.62 -0.08 0.24 0.32 0.03 -0.05 1.63 -0.10 0.01
Half-known vs cipher, equal size 0.52 0.07 0.07 0.07 -0.03 -0.24 3.20 -0.01 0.00
Half-known vs cipher, own steps 0.51 0.13 -0.02 0.01 -0.02 -0.12 3.20 -0.28 0.03
Rank-1 vs cipher, step 10 0.62 -0.05 0.00 -0.01 -0.03 -0.27 3.72 0.00 0.00
Copy vs cipher, step 10 0.61 0.05 -0.08 -0.02 -0.02 0.24 2.12 0.00 0.01
Rank-1 vs cipher, 96 keys 0.56 × × × × 0.20 3.96 -0.08 ×
Copy vs prior-matched cipher, 96 keys 0.39 × × × × 0.70 0.20 -0.10 ×
Half-known vs its shuffle, 96 keys 0.25 × × × × 0.21 2.10 -0.05 ×
Class map vs its shuffled twin -0.04 0.00 × × × -0.08 -0.07 0.00 ×
Case label: real vs memorising twin -0.02 -0.04 × × -0.01 -0.51 0.03 -0.45 -0.07
Repeat label: real vs memorising twin -0.09 -0.01 × × 0.01 -1.09 0.01 -0.20 -0.05
built pairs
real vs shuffled twin
reuse side clearly higher
reuse side higher, one value a side
no clear difference
new side clearly higher
new side higher, one value a side
Each cell: reuse side minus new side (shares and log10 prior: difference of means; size and rank: log10 of the ratio). Seeds per side: 1 to 8. Clearly = seed ranges do not overlap, or the run's filed call.