The re-gating share sets all nine built pairs apart the same way; steering, neurons, size and rank each miss or reverse some, and the prior, one value a side, points the same way on all eight it was read on

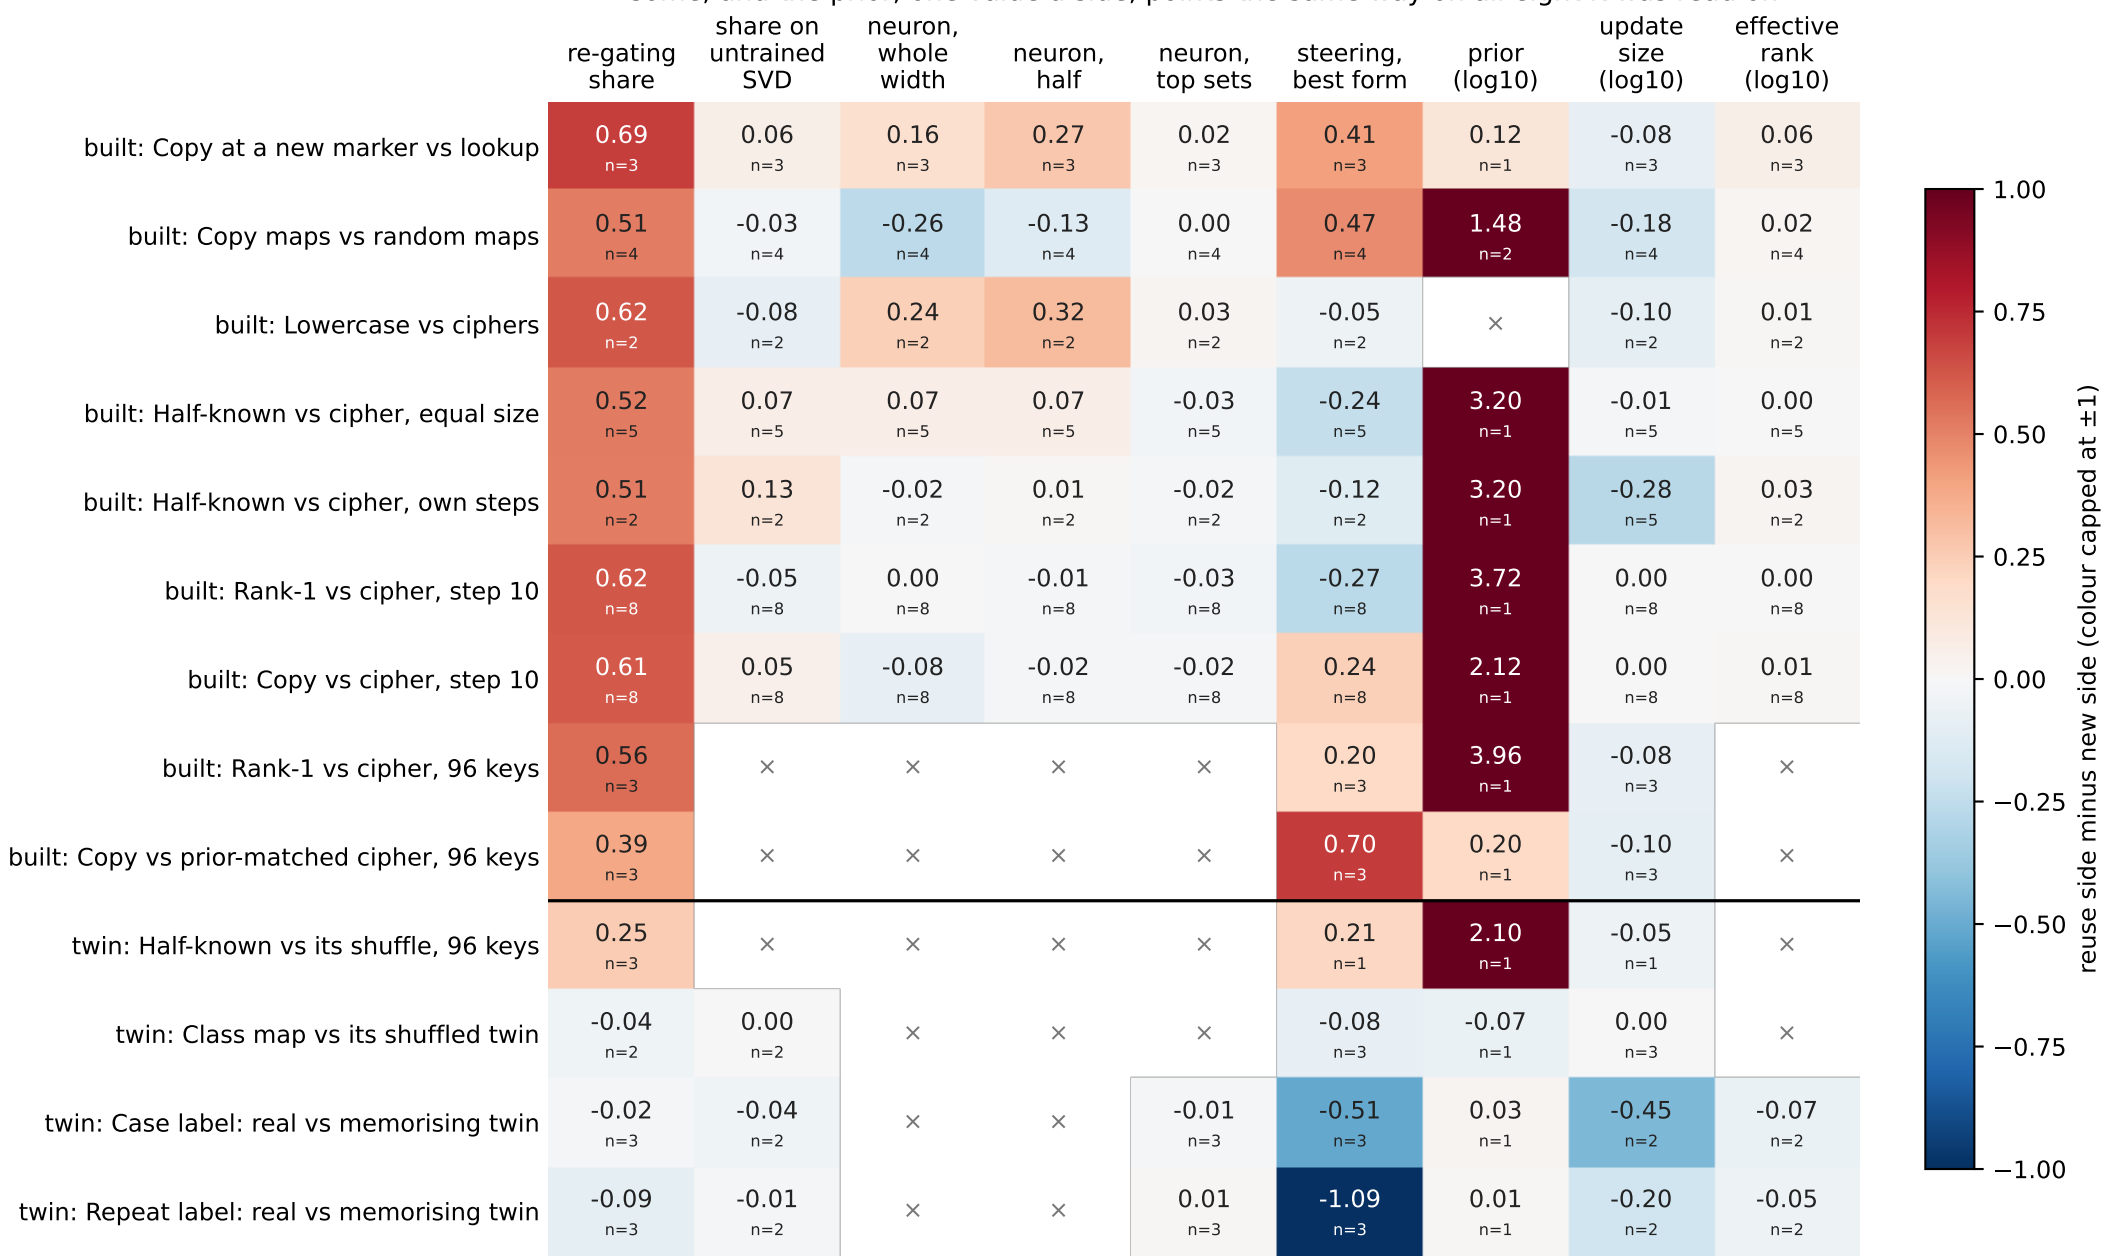

67M model. Reuse side = the side built to reuse, or the real fine-tune; new side = the side built to be new, or the shuffled-label twin. Labels as built, no verdict. Cell = reuse side minus new side, as a difference of side means (shares, log10 prior) or log10 of their ratio (size, rank); n = seeds on the side with fewer; x = no saved reading. Colour is capped at ±1.
re-gating share = part of the fine-tune's gain that appears when only which original subcomponents are switched on changes. share on untrained SVD = the same on each matrix's singular value decomposition, gated by a fixed rule. neuron, whole width = the original MLP weights run on the fine-tune's neuron firing; half = only the MLP input changed; top sets = only each block's top neurons kept. steering = part of the gain from adding the fine-tune's mean activation shift at the read position (pooled, or one vector per key or word; the wider gap is shown). prior = the original model's probability of the trained answer. update size = norm of the weight change. effective rank = how many directions the weight change spreads over.