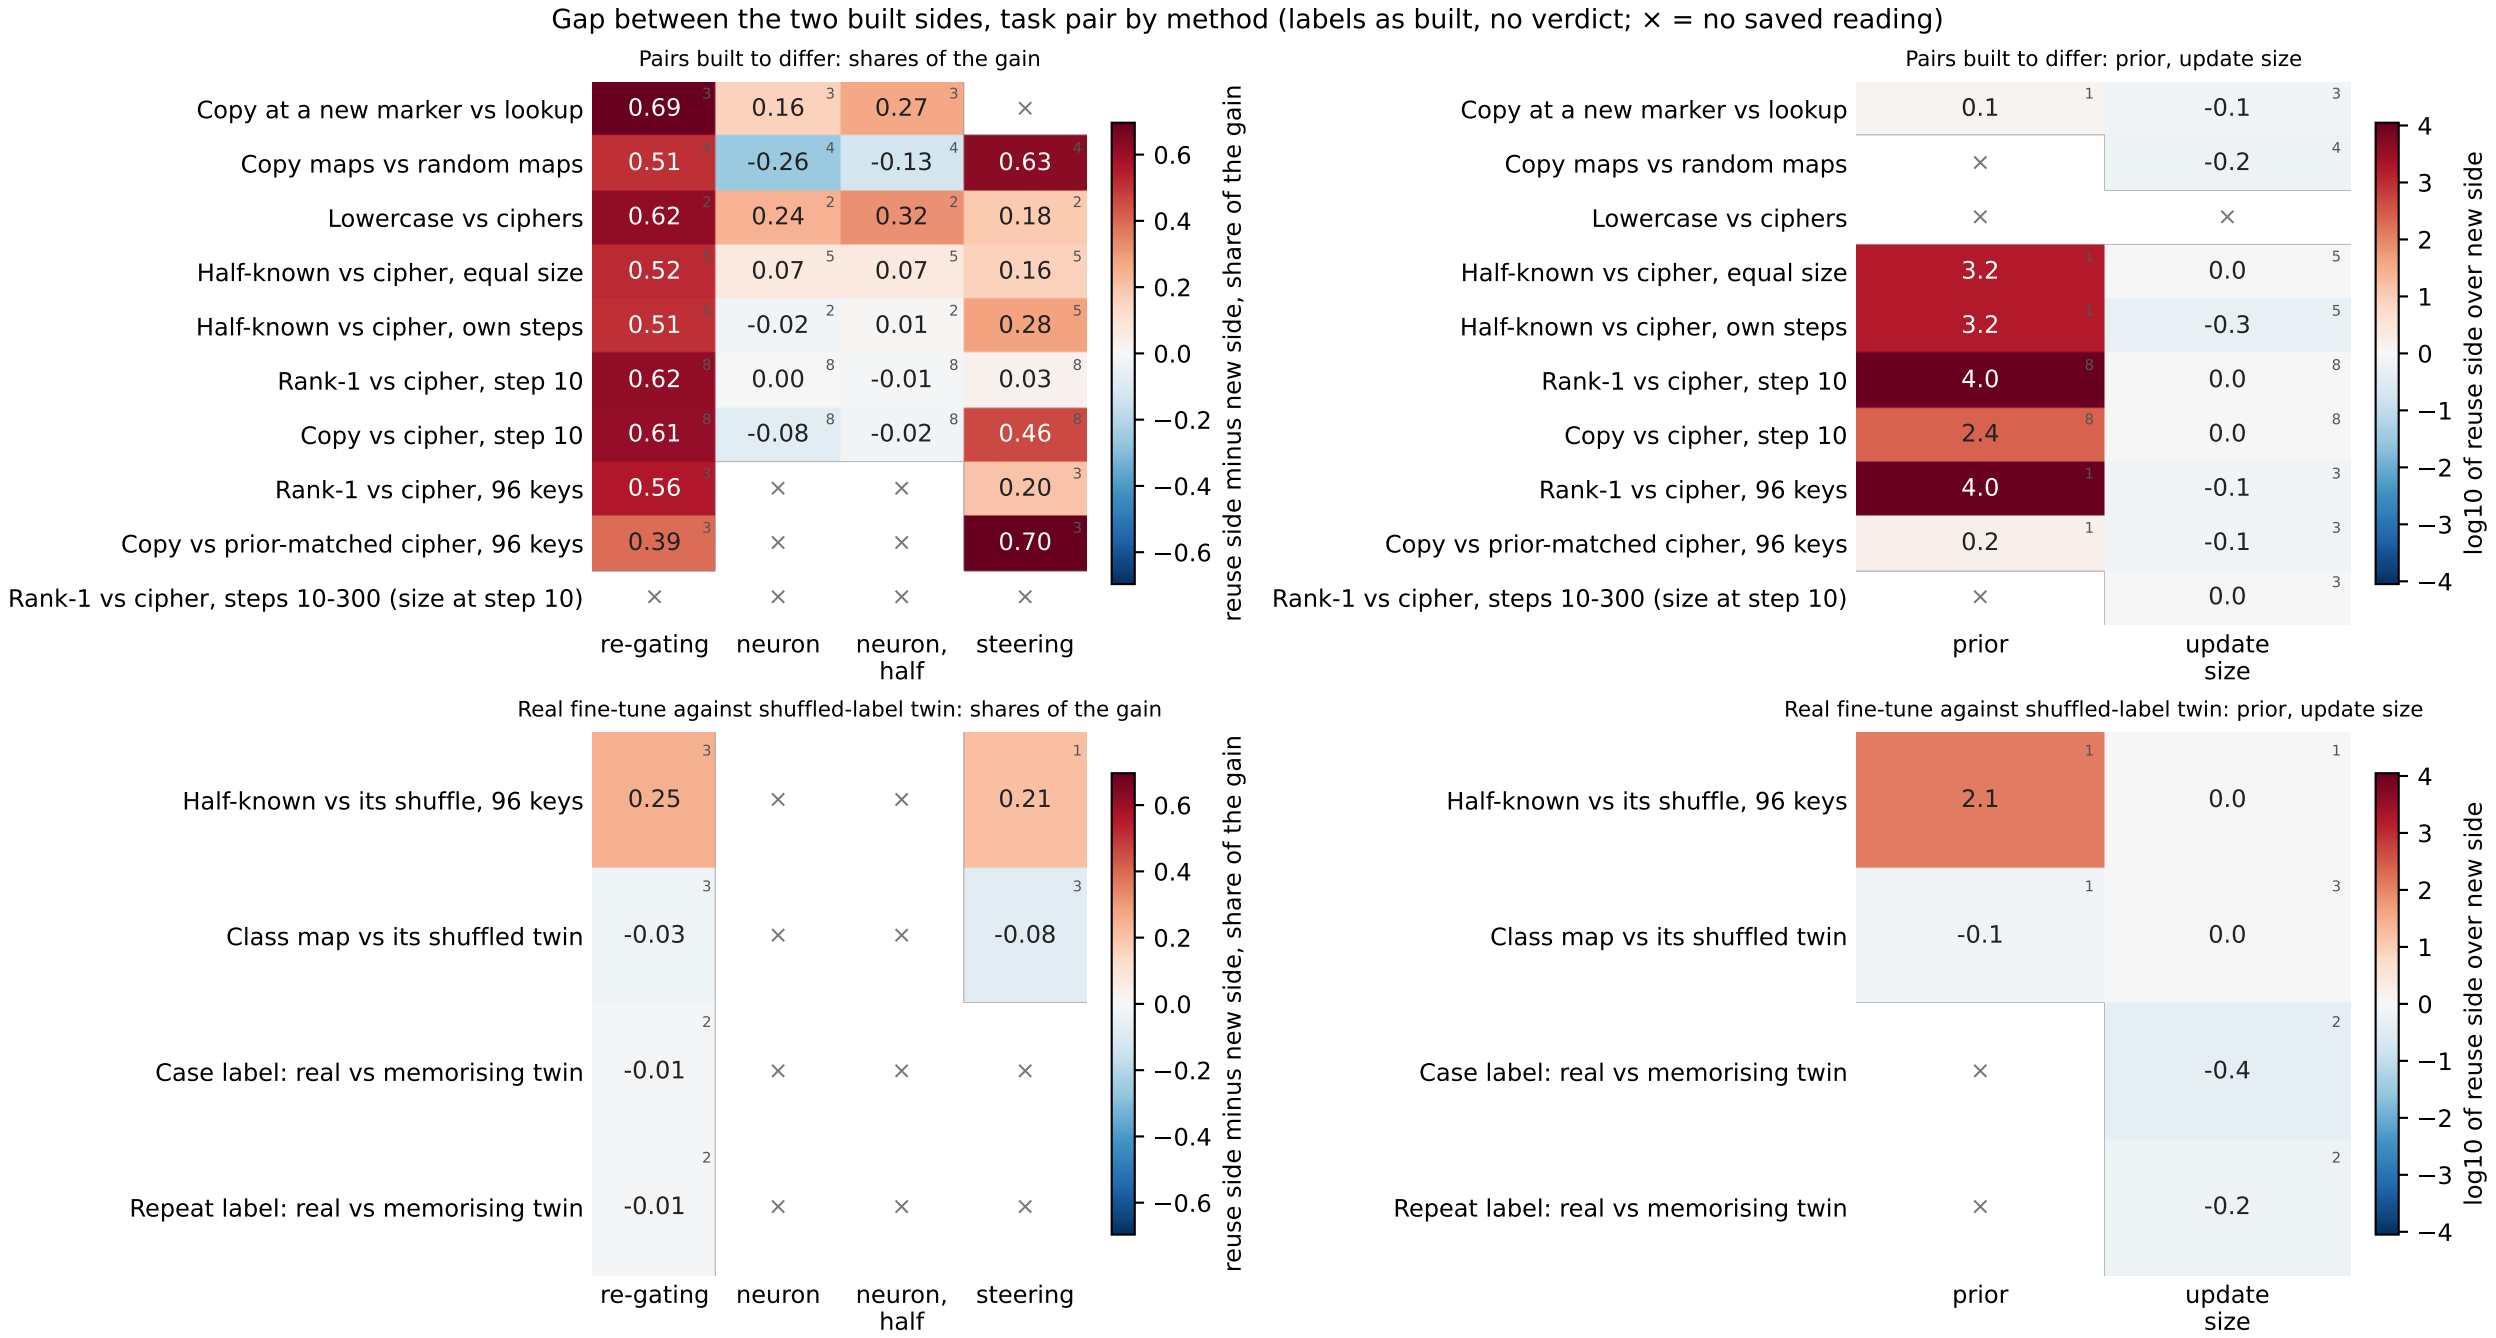

67M model. Reuse side = the first-named task, or the real fine-tune. Gain = log-odds of the trained answer, fine-tuned minus original model. Each cell: difference or log10 ratio of seed means; corner number = seeds on the side with fewer.
re-gating = share of the gain that appears when only which original subcomponents are switched on changes; neuron = the same reading on neuron firing; half = on the MLP input alone;
steering = share of the gain from adding the update's activation shift at the trigger;
prior = the original model's probability of the answer; update size = norm of the weight change. Tasks: reports/analyses/2026-11-09_analysis_scorecard.md.