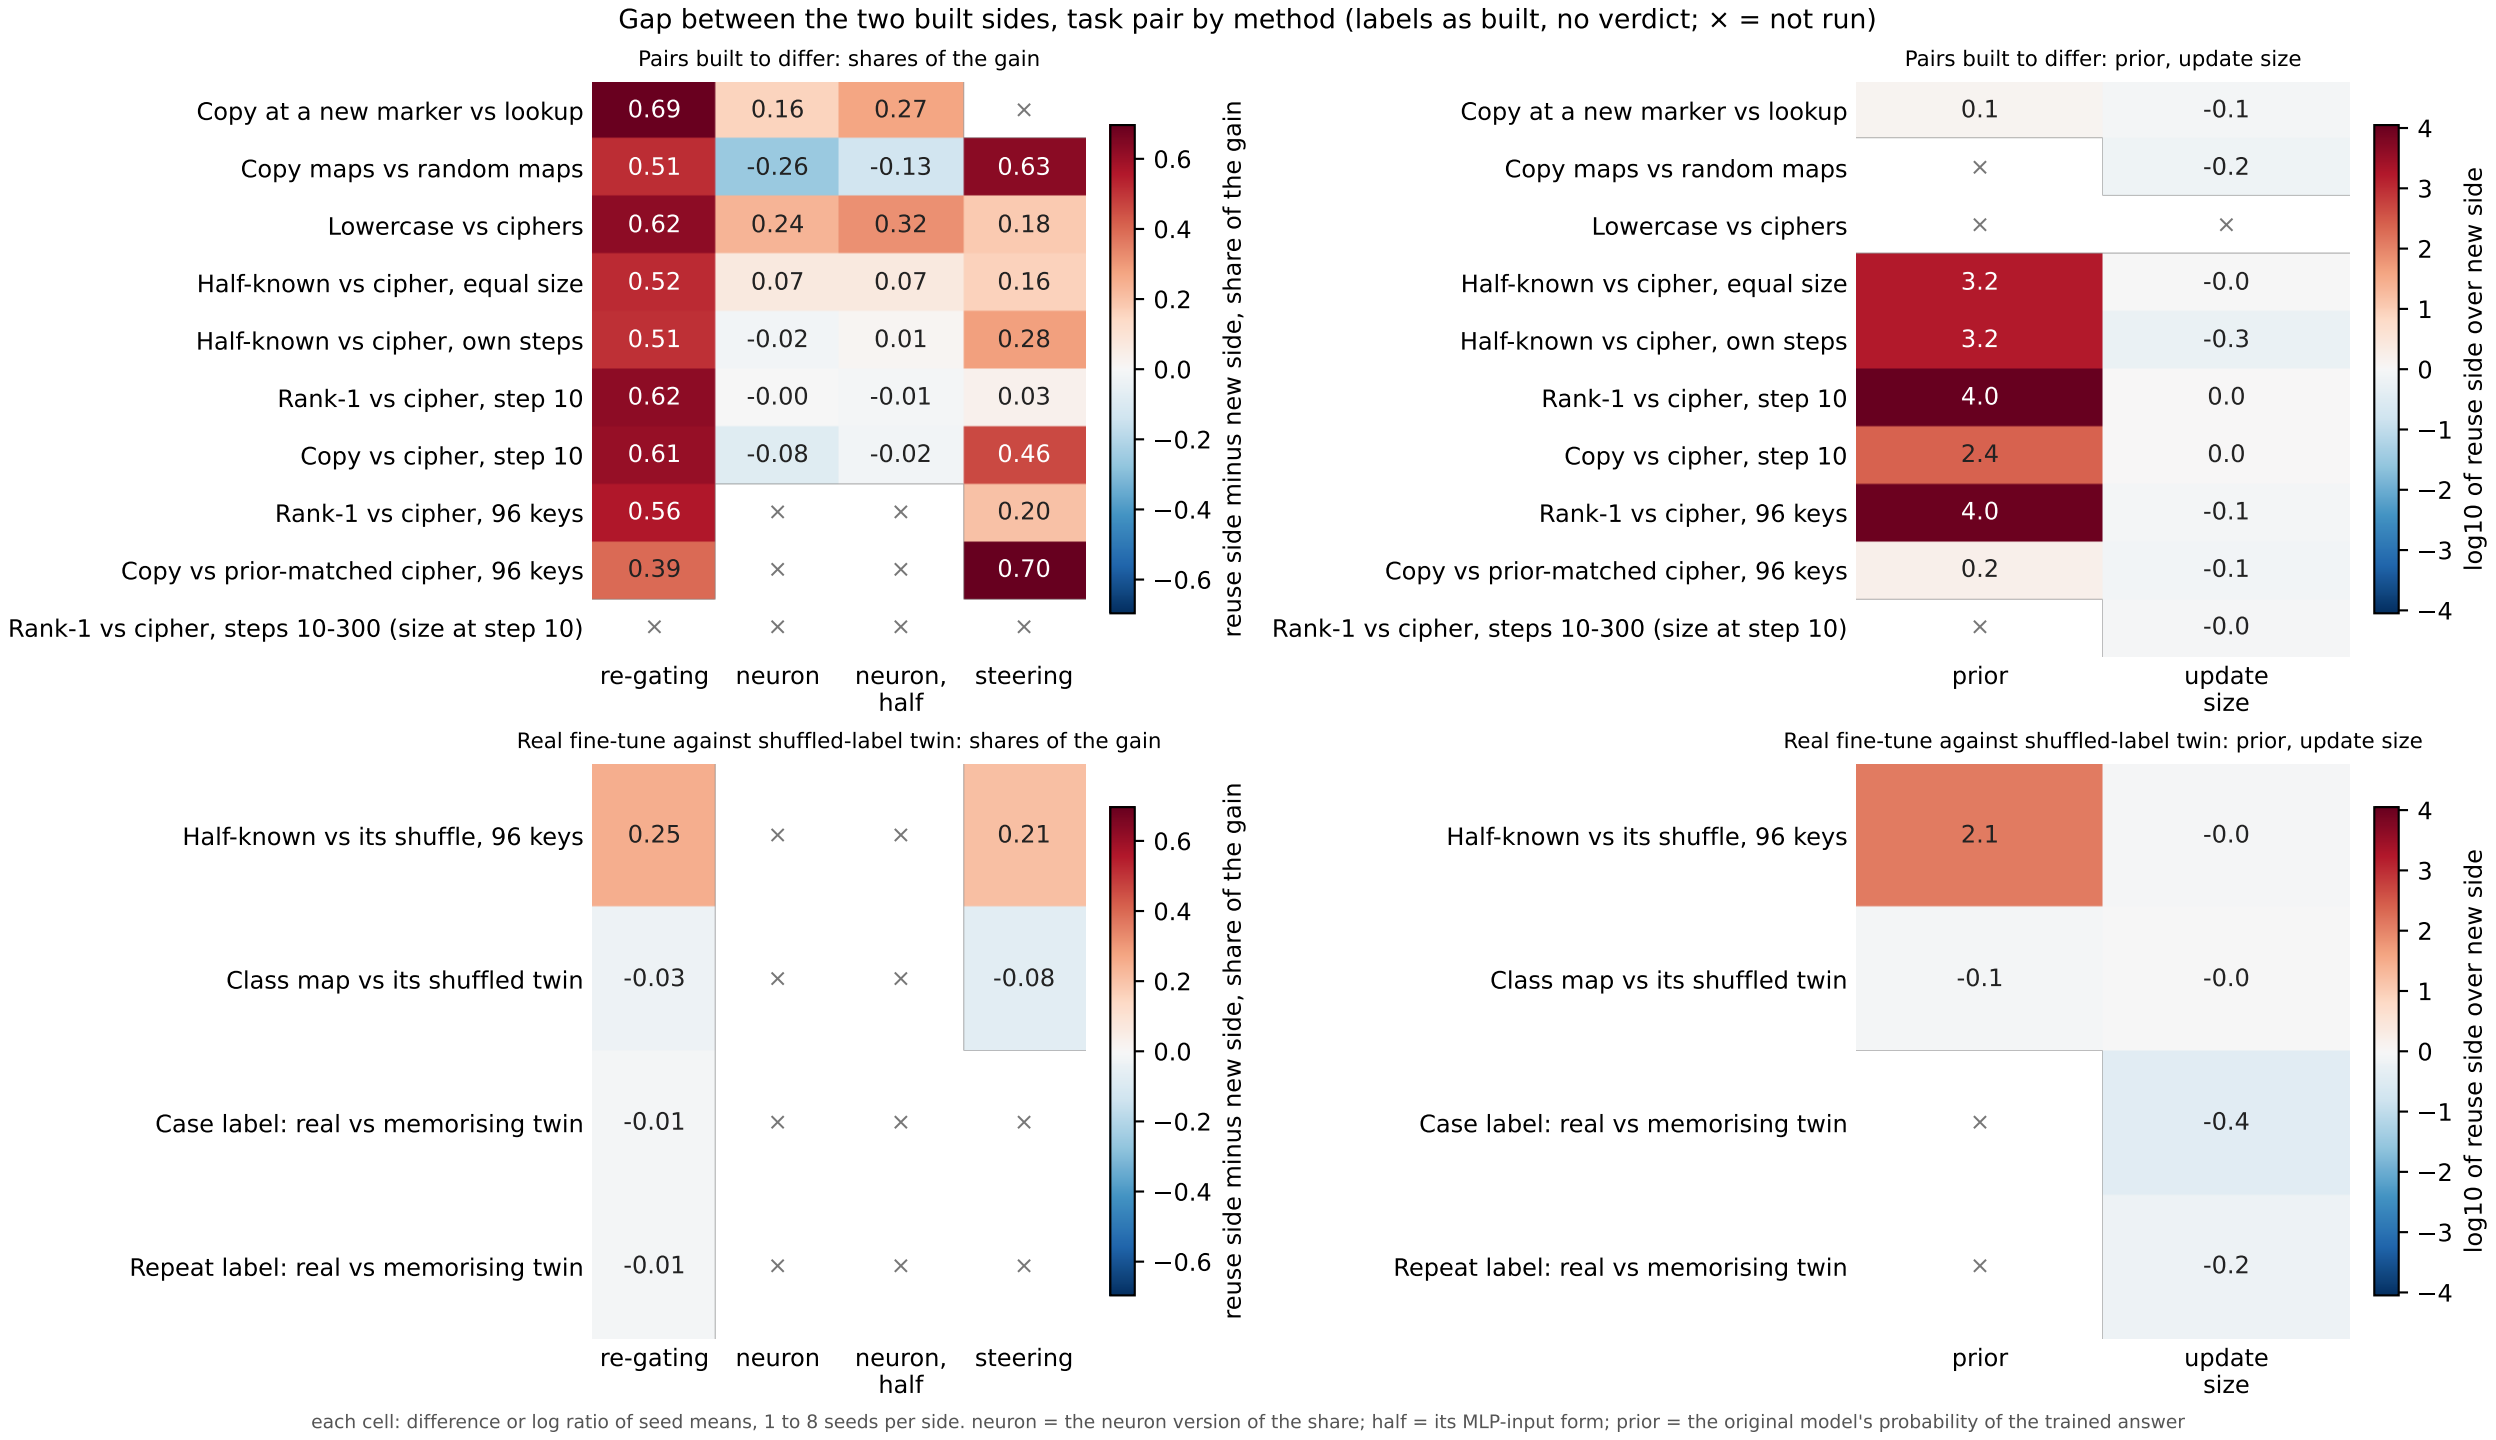
Gap between the two built sides, task pair by method (labels as built, no verdict; × = not run)
Pairs built to differ: shares of the gain
Copy at a new marker vs lookup 0.69 0.16 0.27 ×
Copy maps vs random maps 0.51 -0.26 -0.13 0.63
Lowercase vs ciphers 0.62 0.24 0.32 0.18
Half-known vs cipher, equal size 0.52 0.07 0.07 0.16
Half-known vs cipher, own steps 0.51 -0.02 0.01 0.28
Rank-1 vs cipher, step 10 0.62 -0.00 -0.01 0.03
Copy vs cipher, step 10 0.61 -0.08 -0.02 0.46
Rank-1 vs cipher, 96 keys 0.56 × × 0.20
Copy vs prior-matched cipher, 96 keys 0.39 × × 0.70
Rank-1 vs cipher, steps 10-300 (size at step 10) × × × ×
re-gating
neuron
neuron, half
steering
reuse side minus new side, share of the gain
0.6
0.4
0.2
0.0
−0.2
−0.4
−0.6
Pairs built to differ: prior, update size
Copy at a new marker vs lookup 0.1 -0.1
Copy maps vs random maps × -0.2
Lowercase vs ciphers × ×
Half-known vs cipher, equal size 3.2 -0.0
Half-known vs cipher, own steps 3.2 -0.3
Rank-1 vs cipher, step 10 4.0 0.0
Copy vs cipher, step 10 2.4 0.0
Rank-1 vs cipher, 96 keys 4.0 -0.1
Copy vs prior-matched cipher, 96 keys 0.2 -0.1
Rank-1 vs cipher, steps 10-300 (size at step 10) × -0.0
prior
update size
log10 of reuse side over new side
4
3
2
1
0
−1
−2
−3
−4
Real fine-tune against shuffled-label twin: shares of the gain
Half-known vs its shuffle, 96 keys 0.25 × × 0.21
Class map vs its shuffled twin -0.03 × × -0.08
Case label: real vs memorising twin -0.01 × × ×
Repeat label: real vs memorising twin -0.01 × × ×
re-gating
neuron
neuron, half
steering
reuse side minus new side, share of the gain
0.6
0.4
0.2
0.0
−0.2
−0.4
−0.6
Real fine-tune against shuffled-label twin: prior, update size
Half-known vs its shuffle, 96 keys 2.1 -0.0
Class map vs its shuffled twin -0.1 -0.0
Case label: real vs memorising twin × -0.4
Repeat label: real vs memorising twin × -0.2
prior
update size
log10 of reuse side over new side
4
3
2
1
0
−1
−2
−3
−4
each cell: difference or log ratio of seed means, 1 to 8 seeds per side. neuron = the neuron version of the share; half = its MLP-input form; prior = the original model's probability of the trained answer