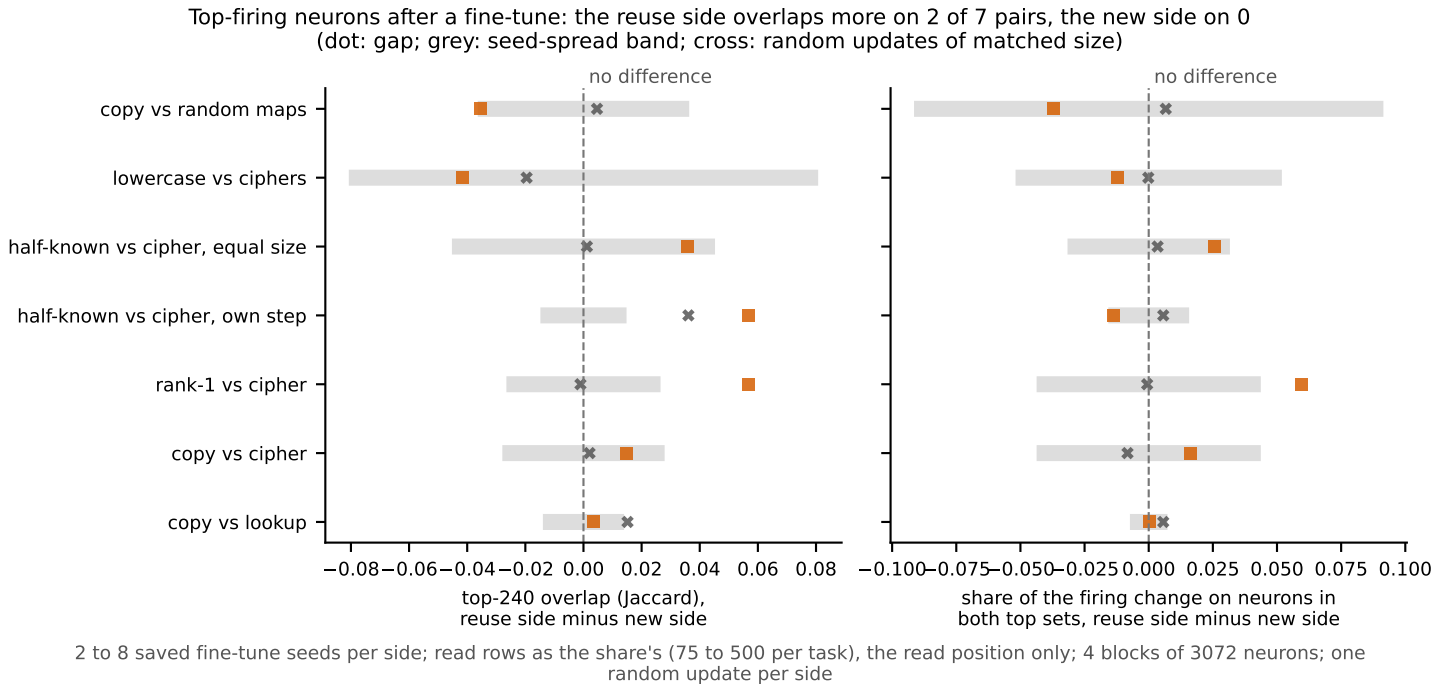
Top-firing neurons after a fine-tune: the reuse side overlaps more on 2 of 7 pairs, the new side on 0
(dot: gap; grey: seed-spread band; cross: random updates of matched size)
no difference
no difference
copy vs random maps
lowercase vs ciphers
half-known vs cipher, equal size
half-known vs cipher, own step
rank-1 vs cipher
copy vs cipher
copy vs lookup
−0.08
−0.06
−0.04
−0.02
0.00
0.02
0.04
0.06
0.08
top-240 overlap (Jaccard),
reuse side minus new side
−0.100
−0.075
−0.050
−0.025
0.000
0.025
0.050
0.075
0.100
share of the firing change on neurons in
both top sets, reuse side minus new side
2 to 8 saved fine-tune seeds per side; read rows as the share's (75 to 500 per task), the read position only; 4 blocks of 3072 neurons; one
random update per side