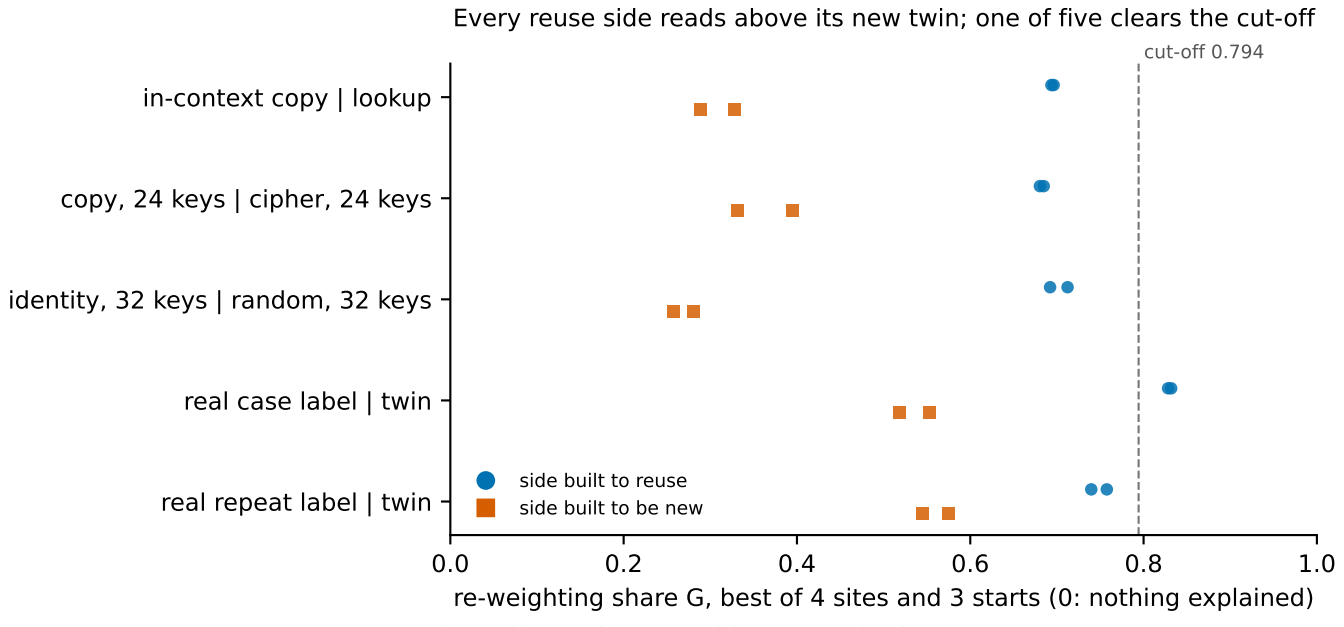
Every reuse side reads above its new twin; one of five clears the cut-off
cut-off 0.794
in-context copy | lookup
copy, 24 keys | cipher, 24 keys
identity, 32 keys | random, 32 keys
real case label | twin
real repeat label | twin
side built to reuse
side built to be new
0.0
0.2
0.4
0.6
0.8
1.0
re-weighting share G, best of 4 sites and 3 starts (0: nothing explained)
2 seeds per side, one dot per saved fine-tune; 8-token inputs