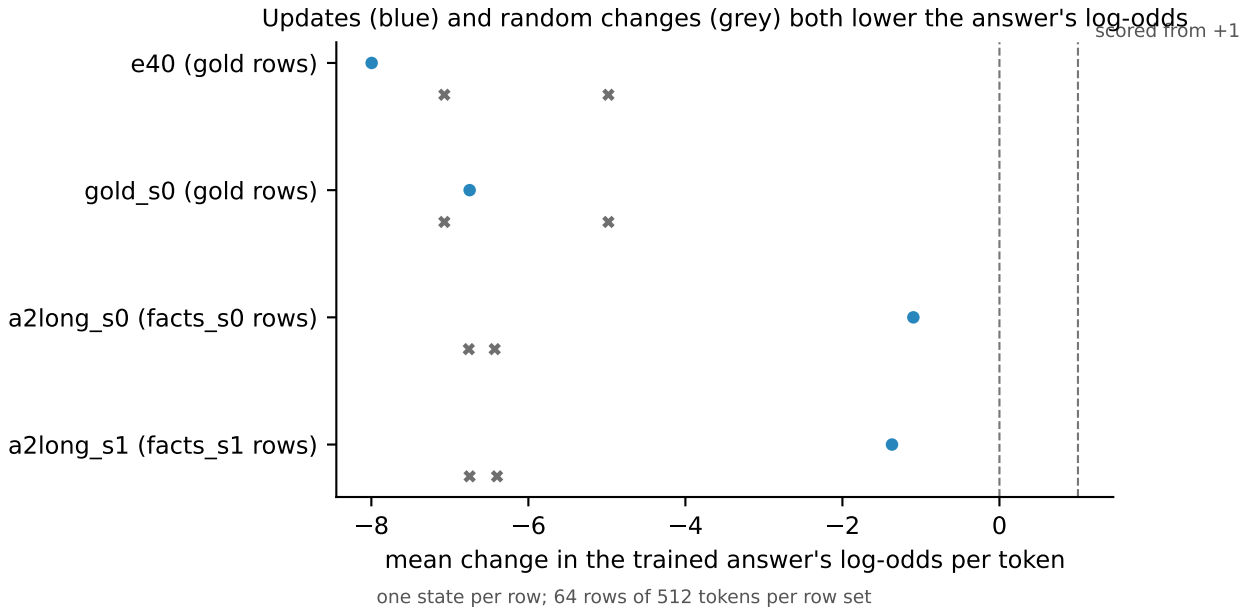
Updates (blue) and random changes (grey) both lower the answer's log-odds
scored from +1
e40 (gold rows)
gold_s0 (gold rows)
a2long_s0 (facts_s0 rows)
a2long_s1 (facts_s1 rows)
−8
−6
−4
−2
0
mean change in the trained answer's log-odds per token
one state per row; 64 rows of 512 tokens per row set