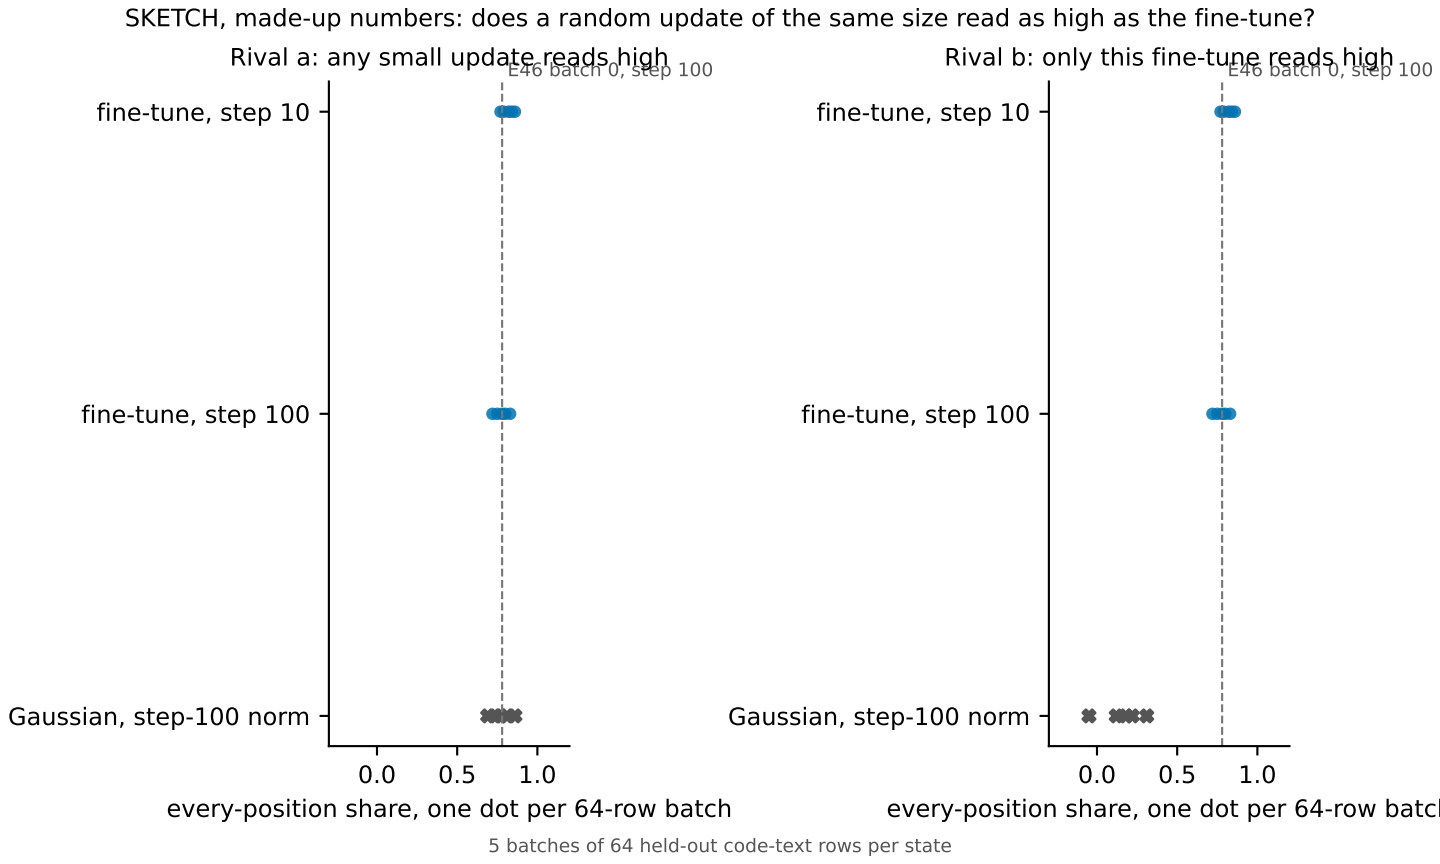
SKETCH, made-up numbers: does a random update of the same size read as high as the fine-tune?
Rival a: any small update reads high
E46 batch 0, step 100
fine-tune, step 10
fine-tune, step 100
Gaussian, step-100 norm
0.0
0.5
1.0
every-position share, one dot per 64-row batch
Rival b: only this fine-tune reads high
E46 batch 0, step 100
fine-tune, step 10
fine-tune, step 100
Gaussian, step-100 norm
0.0
0.5
1.0
every-position share, one dot per 64-row batch
5 batches of 64 held-out code-text rows per state