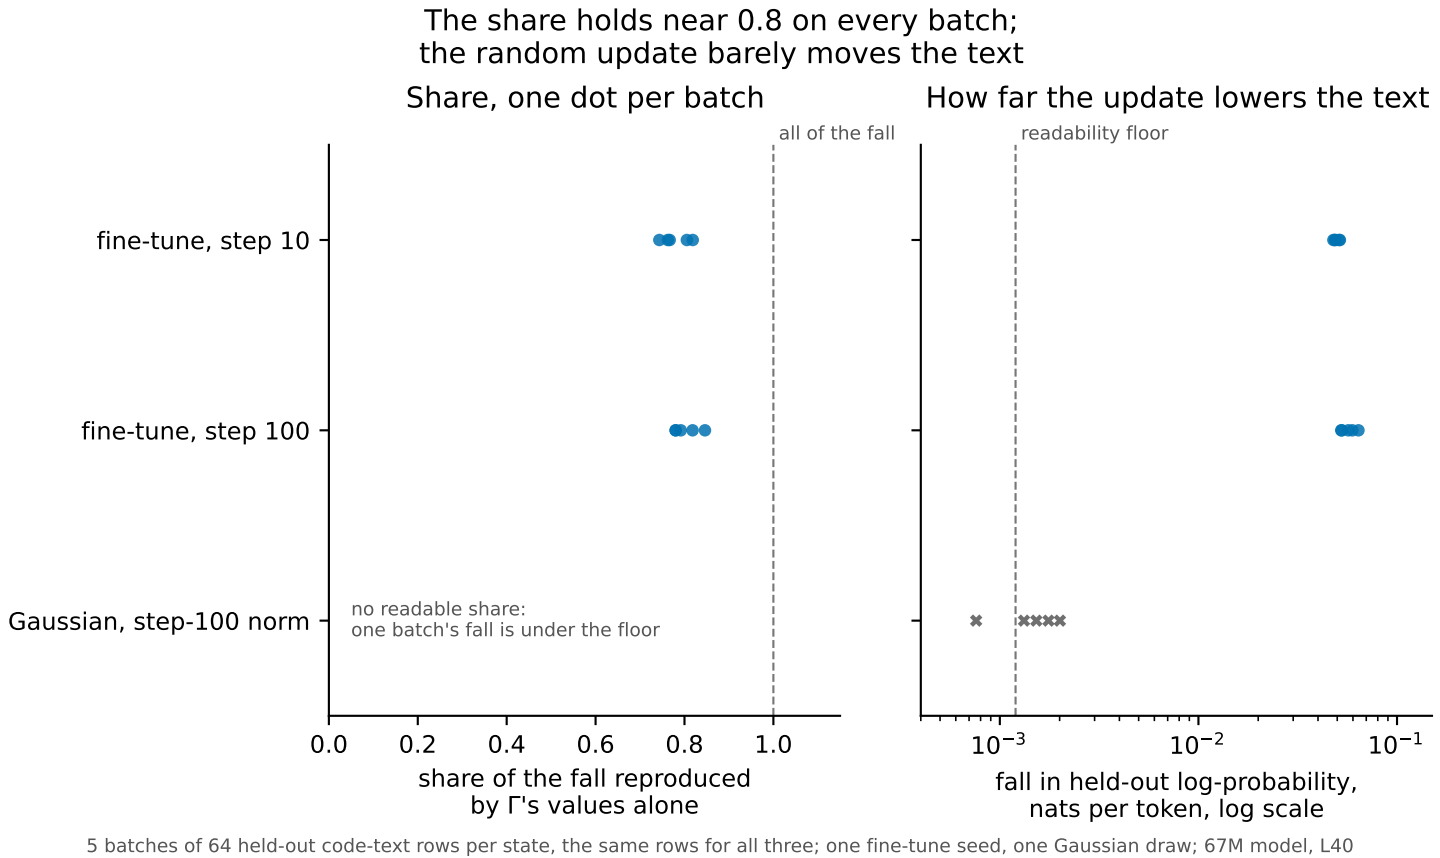
The share holds near 0.8 on every batch;
the random update barely moves the text
Share, one dot per batch
How far the update lowers the text
all of the fall
readability floor
fine-tune, step 10
fine-tune, step 100
Gaussian, step-100 norm
no readable share:
one batch's fall is under the floor
0.0
0.2
0.4
0.6
0.8
1.0
share of the fall reproduced
by Γ's values alone
10^{-3}
10^{-2}
10^{-1}
fall in held-out log-probability,
nats per token, log scale
5 batches of 64 held-out code-text rows per state, the same rows for all three; one fine-tune seed, one Gaussian draw; 67M model, L40